

Project

June 8, 2020

0.1 Imports

```
In [74]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import graphviz
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn import tree
from sklearn import preprocessing
from sklearn.tree import DecisionTreeClassifier
```

0.2 Path

```
In [2]: pwd
```

```
Out[2]: '/Users/ramosem/Documents/SyracuseUniversity/3rd_Quarter/IST707/Project'
```

```
In [3]: path = '/Users/ramosem/Documents/SyracuseUniversity/3rd_Quarter/IST707/Project/'
```

0.3 Read Data

```
In [4]: df = pd.read_csv(path + 'kaggle_Interests_group.csv')
```

```
In [5]: df.head()
```

```
Out[5]:
```

	group	grand_tot_interests	interest1	interest2	interest3	interest4	\
0	C	17	NaN	NaN	NaN	NaN	
1	C	43	1.0	NaN	NaN	NaN	
2	C	27	NaN	NaN	NaN	NaN	
3	C	34	NaN	NaN	NaN	NaN	
4	C	36	NaN	NaN	NaN	NaN	

	interest5	interest6	interest7	interest8	...	interest208	interest209	\
--	-----------	-----------	-----------	-----------	-----	-------------	-------------	---

0	NaN	NaN	NaN	NaN	...	NaN	NaN
1	1.0	NaN	NaN	NaN	...	NaN	NaN
2	NaN	NaN	NaN	NaN	...	NaN	NaN
3	NaN	NaN	NaN	NaN	...	NaN	NaN
4	1.0	NaN	NaN	NaN	...	NaN	NaN

	interest210	interest211	interest212	interest213	interest214	\
0	NaN	NaN	NaN	NaN	NaN	
1	1.0	NaN	NaN	NaN	NaN	
2	1.0	NaN	NaN	NaN	NaN	
3	NaN	1.0	NaN	NaN	NaN	
4	1.0	NaN	NaN	NaN	NaN	

	interest215	interest216	interest217
0	NaN	NaN	NaN
1	1.0	1.0	NaN
2	1.0	1.0	NaN
3	1.0	1.0	NaN
4	1.0	1.0	NaN

[5 rows x 219 columns]

In [6]: `len(df)`

Out[6]: 6340

1 Preprocessing

In [7]: `cat = df['group'].value_counts().reset_index(drop=False)`

In [8]: `cat.columns= ['group', 'count']`

In [9]: `cat['GroupPerc'] = cat['count']/cat['count'].sum()`

In [10]: `cat`

```
Out[10]:
```

	group	count	GroupPerc
0	I	1809	0.285331
1	P	1731	0.273028
2	C	1725	0.272082
3	R	1075	0.169558

In [11]: `df.describe()`

```
Out[11]:
```

	grand_tot_interests	interest1	interest2	interest3	interest4	\
count	6340.000000	993.0	1.0	35.0	25.0	
mean	37.312303	1.0	1.0	1.0	1.0	
std	15.729872	0.0	NaN	0.0	0.0	
min	1.000000	1.0	1.0	1.0	1.0	

25%	28.000000	1.0	1.0	1.0	1.0
50%	39.000000	1.0	1.0	1.0	1.0
75%	48.000000	1.0	1.0	1.0	1.0
max	104.000000	1.0	1.0	1.0	1.0

	interest5	interest6	interest7	interest8	interest9	...	\
count	798.000000	3394.000000	1.0	93.0	333.000000	...	
mean	1.001253	1.000589	1.0	1.0	1.003003	...	
std	0.035400	0.024271	NaN	0.0	0.054800	...	
min	1.000000	1.000000	1.0	1.0	1.000000	...	
25%	1.000000	1.000000	1.0	1.0	1.000000	...	
50%	1.000000	1.000000	1.0	1.0	1.000000	...	
75%	1.000000	1.000000	1.0	1.0	1.000000	...	
max	2.000000	2.000000	1.0	1.0	2.000000	...	

	interest208	interest209	interest210	interest211	interest212	...	\
count	118.000000	110.0	5037.000000	2474.000000	877.000000	...	
mean	1.008475	1.0	1.000596	1.000808	1.002281	...	
std	0.092057	0.0	0.024400	0.028427	0.047727	...	
min	1.000000	1.0	1.000000	1.000000	1.000000	...	
25%	1.000000	1.0	1.000000	1.000000	1.000000	...	
50%	1.000000	1.0	1.000000	1.000000	1.000000	...	
75%	1.000000	1.0	1.000000	1.000000	1.000000	...	
max	2.000000	1.0	2.000000	2.000000	2.000000	...	

	interest213	interest214	interest215	interest216	interest217
count	2.0	72.0	4943.000000	4058.000000	147.0
mean	1.0	1.0	1.000202	1.000246	1.0
std	0.0	0.0	0.014223	0.015698	0.0
min	1.0	1.0	1.000000	1.000000	1.0
25%	1.0	1.0	1.000000	1.000000	1.0
50%	1.0	1.0	1.000000	1.000000	1.0
75%	1.0	1.0	1.000000	1.000000	1.0
max	1.0	1.0	2.000000	2.000000	1.0

[8 rows x 218 columns]

Columns

```
In [12]: yParm = ['group']
        yParmStr = 'group'
        xParm = df.columns[2:]
```

1.0.1 Get rid of the ones with a small and big amounts of positive responses

```
In [13]: goodCols = []
        for col in xParm:
            if (df[col].count() < 10):
```

```

        #         print(col)
        continue
    elif (df[col].count() > 5990):
        print(col)
        continue
    else:
        goodCols.append(col)

```

interest162

interest183

In [14]: len(goodCols)

Out[14]: 175

In [15]: len(xParm) - len(goodCols)

Out[15]: 42

In [16]: xParm = goodCols

1.0.2 Change twos to ones

```

In [17]: for col in xParm:
        df[col] = [x if pd.isnull(x) else 1 for x in df[col].values]

```

1.0.3 Get rid off of well distributed interests among all categories.

In [18]: catCount = df[xParm + yParm].groupby(yParmStr).count().transpose()

In [19]: catCount['Total'] = catCount[['C', 'I', 'P', 'R']].sum(axis=1)

```

In [20]: for col in ['C', 'I', 'P', 'R']:
        catCount[col+'_perc'] = catCount[col]*100/cat.loc[cat['group']==col]['count'].val

```

In [21]: catCount['TotalPerc'] = catCount[['C_perc', 'I_perc', 'P_perc', 'R_perc']].sum(axis=1)

```

In [22]: for col in ['C', 'I', 'P', 'R']:
        catCount[col+'_norm'] = catCount[col+'_perc']*100/catCount['TotalPerc']

```

```

In [23]: badRow = []
        goodRow = []
        for idx, row in catCount.iterrows():
            if (row['C_norm'] > 30) | (row['C_norm'] < 20):
                goodRow.append(idx)
                continue
            elif (row['I_norm'] > 30) | (row['I_norm'] < 20):
                goodRow.append(idx)
                continue

```

```

elif (row['P_norm'] > 30) | (row['P_norm'] < 20):
    goodRow.append(idx)
    continue
elif (row['R_norm'] > 30) | (row['R_norm'] < 20):
    goodRow.append(idx)
    continue
else:
    badRow.append(idx)

```

In [24]: `len(badRow)`

Out[24]: 51

In [25]: `catCount.loc[catCount.index.isin(badRow)].head()`

```

Out[25]: group      C      I      P      R  Total      C_perc      I_perc      P_perc  \
interest6      859    1032     888    615    3394  49.797101  57.048093  51.299827
interest12    1211    1235    1248    825    4519  70.202899  68.269762  72.097054
interest14      86     111     100     68     365   4.985507   6.135987   5.777008
interest15     455     545     559    333    1892  26.376812  30.127142  32.293472
interest16    1218    1219    1251    823    4511  70.608696  67.385296  72.270364

      group      R_perc  TotalPerc      C_norm      I_norm      P_norm      R_norm
interest6  57.209302  215.354323  23.123335  26.490340  23.821127  26.565198
interest12  76.744186  287.313901  24.434216  23.761385  25.093479  26.710920
interest14   6.325581   23.224083  21.466971  26.420792  24.875073  27.237163
interest15  30.976744  119.774170  22.022120  25.153288  26.961967  25.862625
interest16  76.558140  286.822495  24.617559  23.493728  25.196895  26.691818

```

In [26]: `xParm = goodRow`

In [27]: `len(xParm)`

Out[27]: 124

1.0.4 Get rid of Anomalous Individuals

In [28]: `df['grand_tot_interests'] = df[xParm].sum(axis=1)`

Zero Interests

In [29]: `df = df.loc[df['grand_tot_interests']!=0].reset_index(drop=True)`

In [30]: `len(df)`

Out[30]: 5788


Anomaly in distribution

```
In [31]: df['grand_tot_interests'].describe()
```

```
Out[31]: count      5788.000000  
         mean         9.720974  
         std         6.236600  
         min          1.000000  
         25%          5.000000  
         50%          9.000000  
         75%         14.000000  
         max         52.000000  
         Name: grand_tot_interests, dtype: float64
```

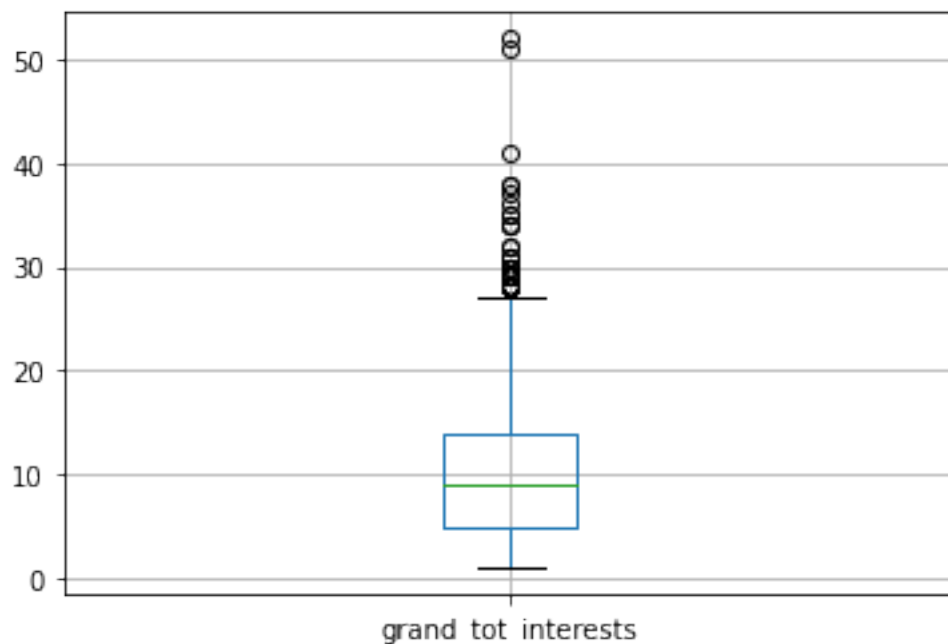
```
In [32]: df['grand_tot_interests'].hist(bins=40)
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x102865358>
```



```
In [33]: df[['grand_tot_interests']].boxplot()
```

```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0x1166189b0>
```



```
In [34]: df['interest9'].value_counts()
```

```
Out[34]: 1.0      333
         Name: interest9, dtype: int64
```

```
In [35]: df[xParm].describe()
```

```
Out[35]:
```

	interest1	interest3	interest4	interest5	interest8	interest9	\
count	993.0	35.0	25.0	798.0	93.0	333.0	
mean	1.0	1.0	1.0	1.0	1.0	1.0	
std	0.0	0.0	0.0	0.0	0.0	0.0	
min	1.0	1.0	1.0	1.0	1.0	1.0	
25%	1.0	1.0	1.0	1.0	1.0	1.0	
50%	1.0	1.0	1.0	1.0	1.0	1.0	
75%	1.0	1.0	1.0	1.0	1.0	1.0	
max	1.0	1.0	1.0	1.0	1.0	1.0	

	interest11	interest13	interest19	interest20	...	interest201	\
count	175.0	18.0	141.0	165.0	...	1671.0	
mean	1.0	1.0	1.0	1.0	...	1.0	
std	0.0	0.0	0.0	0.0	...	0.0	
min	1.0	1.0	1.0	1.0	...	1.0	
25%	1.0	1.0	1.0	1.0	...	1.0	
50%	1.0	1.0	1.0	1.0	...	1.0	
75%	1.0	1.0	1.0	1.0	...	1.0	
max	1.0	1.0	1.0	1.0	...	1.0	

	interest203	interest204	interest205	interest208	interest209 \
count	26.0	117.0	134.0	118.0	110.0
mean	1.0	1.0	1.0	1.0	1.0
std	0.0	0.0	0.0	0.0	0.0
min	1.0	1.0	1.0	1.0	1.0
25%	1.0	1.0	1.0	1.0	1.0
50%	1.0	1.0	1.0	1.0	1.0
75%	1.0	1.0	1.0	1.0	1.0
max	1.0	1.0	1.0	1.0	1.0


	interest211	interest212	interest214	interest217
count	2474.0	877.0	72.0	147.0
mean	1.0	1.0	1.0	1.0
std	0.0	0.0	0.0	0.0
min	1.0	1.0	1.0	1.0
25%	1.0	1.0	1.0	1.0
50%	1.0	1.0	1.0	1.0
75%	1.0	1.0	1.0	1.0
max	1.0	1.0	1.0	1.0

[8 rows x 124 columns]

```
In [36]: df['Zscore'] = df.groupby('group')['grand_tot_interests'].apply(lambda x: x.div(x.mean()))
```

```
plt.figure(figsize=(7,5))
df.groupby('group').Zscore.plot.kde()
plt.legend()
```

```
Out[36]: <matplotlib.legend.Legend at 0x116724a58>
```





```
In [37]: df['grand_tot_interests']
```

```
Out[37]: 0      15.0
         1       9.0
         2       7.0
         3      18.0
         4       4.0
         ...
        5783    14.0
        5784     9.0
        5785    12.0
        5786    13.0
        5787    26.0
        Name: grand_tot_interests, Length: 5788, dtype: float64
```

```
In [38]: df.boxplot(column='grand_tot_interests',by='group')
```

```
Out[38]: <matplotlib.axes._subplots.AxesSubplot at 0x11680dd30>
```



```
In [39]: df = df.loc[df['grand_tot_interests']<28].reset_index(drop=True)
```

```
In [40]: len(df)
```

```
Out[40]: 5737
```

2 Analysis

```
In [45]: for col in xParm:
          df[col] = [0 if pd.isnull(x) else 1 for x in df[col].values]
```

2.1 Decision Tree

```
In [46]: parameters = {'n_estimators':[10, 30, 50, 70], 'criterion':['entropy'],
                       'max_depth':[10, 20, 25]}
```

```
In [47]: rf = GridSearchCV(RandomForestClassifier(), parameters, n_jobs=5, cv=5, verbose=10)
```

```
In [48]: rf = rf.fit(df[xParm].values, df[yParm].astype(str).values)
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

```
[Parallel(n_jobs=5)]: Using backend LokyBackend with 5 concurrent workers.
[Parallel(n_jobs=5)]: Done   3 tasks      | elapsed:    0.2s
[Parallel(n_jobs=5)]: Done   8 tasks      | elapsed:    0.6s
[Parallel(n_jobs=5)]: Done  15 tasks      | elapsed:    1.2s
[Parallel(n_jobs=5)]: Done  22 tasks      | elapsed:    2.2s
[Parallel(n_jobs=5)]: Done  31 tasks      | elapsed:    3.8s
[Parallel(n_jobs=5)]: Done  40 tasks      | elapsed:    5.0s
[Parallel(n_jobs=5)]: Done  51 tasks      | elapsed:    6.4s
[Parallel(n_jobs=5)]: Done  58 out of  60 | elapsed:    7.7s remaining:    0.3s
[Parallel(n_jobs=5)]: Done  60 out of  60 | elapsed:    7.7s finished
/Users/ramosem/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_search.py:739: D
self.best_estimator_.fit(X, y, **fit_params)
```

2.1.1 Check Out the Best Model

```
In [49]: rf_model = rf.best_estimator_
         print (rf.best_score_, rf.best_params_)
```

```
0.5847999173726948 {'criterion': 'entropy', 'max_depth': 25, 'n_estimators': 70}
```

```
In [50]: rf_model
```

```
Out[50]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                                criterion='entropy', max_depth=25, max_features='auto',
                                max_leaf_nodes=None, max_samples=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=70,
                                n_jobs=None, oob_score=False, random_state=None,
                                verbose=0, warm_start=False)
```

2.1.2 Prediction

```
In [51]: df['predict_rf'] = rf_model.predict(df[xParm].values)
```

```
In [52]: df[['predict_rf', 'group']].head()
```

```
Out[52]:   predict_rf group
0           C      C
1           C      C
2           C      C
3           C      C
4           C      C
```

```
In [53]: df['correct_rf'] = df['predict_rf'] == df['group']
```

```
In [54]: df[['group', 'correct_rf']].groupby(['group', 'correct_rf']).size().unstack(fill_valu
```

```
Out[54]: correct_rf group  False  True
         0             C    156  1364
         1             I     97  1588
         2             P     45  1531
         3             R    144   812
```

```
In [55]: df['correct_rf'].value_counts()
```

```
Out[55]: True      5295
         False     442
         Name: correct_rf, dtype: int64
```

```
In [56]: print("Total Correctly Predicted ", len(df.loc[df['correct_rf']])/len(df))
```

```
Total Correctly Predicted  0.9229562489105805
```

2.1.3 Find Splits

```
In [57]: colImp = pd.DataFrame(zip(xParm, rf_model.feature_importances_), columns=['Interest',
```

```
In [116]: xParm = colImp.sort_values('Importance', ascending=False)['Interest'].values[0:30]
```

3 Decision Tree

3.1 Train

```
In [117]: parameters = {'criterion':['entropy'],
                        'max_depth':[5,10,15,20], 'max_features':[5,10,15,20, 30], 'max_leaf_no
```

```
In [118]: dt = GridSearchCV(DecisionTreeClassifier(), parameters, n_jobs=5, cv=5, verbose=10)
```

```
In [119]: dt = dt.fit(df[xParm].values, df[yParm].values)
```

```
Fitting 5 folds for each of 60 candidates, totalling 300 fits
```

```
[Parallel(n_jobs=5)]: Using backend LokyBackend with 5 concurrent workers.
[Parallel(n_jobs=5)]: Batch computation too fast (0.0414s.) Setting batch_size=2.
[Parallel(n_jobs=5)]: Done   3 tasks      | elapsed:    0.1s
[Parallel(n_jobs=5)]: Done   8 tasks      | elapsed:    0.1s
[Parallel(n_jobs=5)]: Batch computation too fast (0.0584s.) Setting batch_size=4.
[Parallel(n_jobs=5)]: Done  20 tasks      | elapsed:    0.1s
[Parallel(n_jobs=5)]: Batch computation too fast (0.1440s.) Setting batch_size=8.
[Parallel(n_jobs=5)]: Done  38 tasks      | elapsed:    0.3s
[Parallel(n_jobs=5)]: Done  78 tasks      | elapsed:    0.6s
[Parallel(n_jobs=5)]: Done 150 tasks      | elapsed:    0.9s
[Parallel(n_jobs=5)]: Done 238 tasks      | elapsed:    1.6s
[Parallel(n_jobs=5)]: Done 263 tasks      | elapsed:    1.8s
[Parallel(n_jobs=5)]: Done 290 tasks      | elapsed:    1.9s
[Parallel(n_jobs=5)]: Done 300 out of 300 | elapsed:    2.0s finished
```

3.2 Check out Best Model

```
In [120]: dt_model = dt.best_estimator_  
          print (dt.best_score_, dt.best_params_)
```

```
0.5600478752327691 {'criterion': 'entropy', 'max_depth': 20, 'max_features': 15, 'max_leaf_nodes': 10}
```

```
In [121]: dt_model
```

```
Out[121]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',  
                                max_depth=20, max_features=15, max_leaf_nodes=10,  
                                min_impurity_decrease=0.0, min_impurity_split=None,  
                                min_samples_leaf=1, min_samples_split=2,  
                                min_weight_fraction_leaf=0.0, presort='deprecated',  
                                random_state=None, splitter='best')
```

3.3 Predict

```
In [122]: df['predict_dt'] = dt_model.predict(df[xParm].values)
```

```
In [123]: df[['predict_dt', 'group']].head()
```

```
Out[123]:   predict_dt group  
0           C      C  
1           I      C  
2           C      C  
3           C      C  
4           C      C
```

```
In [124]: df['correct_dt'] = df['predict_dt'] == df['group']
```

```
In [125]: df[['group', 'correct_dt']].groupby(['group', 'correct_dt']).size().unstack(fill_value=0)
```

```
Out[125]: correct_dt group  False  True  
0           C      690    830  
1           I      561   1124  
2           P      388   1188  
3           R      849    107
```

```
In [126]: df['correct_dt'].value_counts()
```

```
Out[126]: True      3249  
         False    2488  
         Name: correct_dt, dtype: int64
```

```
In [127]: print("Total Correctly Predicted ", len(df.loc[df['correct_dt']])/len(df))
```

```
Total Correctly Predicted  0.5663238626459822
```

3.4 Plot Decision Tree

```
In [128]: dot_data = tree.export_graphviz(dt_model, out_file=None)
graph = graphviz.Source(dot_data)
graph.render("Personality Decision Tree")
```

```
Out[128]: 'Personality Decision Tree.pdf'
```

```
In [129]: dot_data = tree.export_graphviz(dt_model, out_file=None,
feature_names=xParm,
class_names=df['group'].unique(),
filled=True, rounded=True,
special_characters=True)

graph = graphviz.Source(dot_data)
```

```
In [153]: # plt.figure(figsize=(5,5))
# graph
```

4 Group By Group in Remaining Parameters

```
In [150]: df[np.append(xParm, yParm)].groupby(yParmStr).sum().transpose().reset_index(drop=False)
```

```
Out[150]:
```

group	index	C	I	P	R
0	interest201	300	1131	94	112
1	interest200	892	1455	432	497
2	interest48	174	935	646	100
3	interest196	550	48	42	42
4	interest135	560	122	137	66
5	interest211	471	792	835	330
6	interest79	548	780	493	332
7	interest63	778	1143	773	541
8	interest112	631	803	571	368
9	interest164	951	1448	1191	583

```
In [151]: df[np.append(xParm, yParm)].groupby(yParmStr).sum().transpose().reset_index(drop=False)
```

```
Out[151]:
```

group	index	C	I	P	R
10	interest62	738	1030	626	444
11	interest186	62	225	257	41
12	interest147	278	394	287	181
13	interest44	293	481	353	195
14	interest65	336	383	268	237
15	interest88	185	100	116	195
16	interest187	216	224	205	169
17	interest1	269	324	197	179
18	interest149	428	578	360	245
19	interest181	190	348	184	104

```
In [152]: df[np.append(xParm, yParm)].groupby(yParmStr).sum().transpose().reset_index(drop=False)
```

```

Out[152]: group      index      C      I      P      R
20      interest142  424    574    386    258
21      interest144  347    523    342    216
22      interest212  181    301    238    116
23      interest43   217    321    229    144
24      interest82   286    161    160     92
25      interest153  133     91     62    128
26      interest41   194    182    118    130
27      interest146  139     93    117     84
28      interest118  127    199    184    100
29      interest5    296    213    183     76

```